



✓ **Parabéns! Você foi aprovado!**
PARA SER APROVADO 80% ou superior

Continue aprendendo

NOTA
100%

MDPs

NÚMERO TOTAL DE PONTOS 12

1. The learner and decision maker is the ____.

1 / 1 ponto

- ☐ State
- ☐ Environment
- ☐ Reward
- ☒ Agent

✓ **Correto**
Correct!

2. At each time step the agent takes an ____.

1 / 1 ponto

- ☐ Reward
- ☒ Action
- ☐ State
- ☐ Environment

✓ **Correto**
Correct!

3. Imagine the agent is learning in an episodic problem. Which of the following is true?

1 / 1 ponto

- ☐ The number of steps in an episode is always the same.
- ☐ The agent takes the same action at each step during an episode.
- ☒ The number of steps in an episode is stochastic: each episode can have a different number of steps.

✓ **Correto**
Correct!

4. If the reward is always +1 what is the sum of the discounted infinite return when $\gamma < 1$

1 / 1 ponto

$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$

- ☐ Infinity.
- ☐ $G_t = \frac{\gamma}{1-\gamma}$
- ☐ $G_t = 1 + \gamma^k$
- ☒ $G_t = \frac{1}{1-\gamma}$

✓ **Correto**
Correct!

5. What is the difference between a small gamma (discount factor) and a large gamma?

1 / 1 ponto

- ☒ With a larger discount factor the agent is more far-sighted and considers rewards farther into the future.
- ☐ With a smaller discount factor the agent is more far-sighted and considers rewards farther into the future.
- ☐ The size of the discount factor has no effect on the agent.

✓ **Correto**
Correct!

6. Suppose $\gamma = 0.8$ and the reward sequence is $R_1 = 5$ followed by an infinite sequence of 10s. What is G_0 ?

1 / 1 ponto

- ☒ 45
- ☐ 55
- ☐ 15

✓ **Correto**
Correct!

$G_2 = 10 / (1 - 0.8) = 50$

$G_1 = 10 + .8 * (50) = 50$

$G_0 = 5 + .8 * 50 = 45$

7. What does MDP stand for?

1 / 1 ponto

- ☐ Meaningful Decision Process
- ☐ Markov Decision Process

- ☐ Markov Decision Protocol
- ☐ Markov Deterministic Policy
- ☒ Markov Decision Process

✓ **Correto**
Correct!

8. Consider using reinforcement learning to control the motion of a robot arm in a repetitive pick-and-place task. If we want to learn movements that are fast and smooth, the learning agent will have to control the motors directly and have low-latency information about the current positions and velocities of the mechanical linkages. The actions in this case might be the voltages applied to each motor at each joint, and the states might be the latest readings of joint angles and velocities. The reward might be +1 for each object successfully picked up and placed. To encourage smooth movements, on each time step a small, negative reward can be given as a function of the moment-to-moment "jerkiness" of the motion. Is this a valid MDP?

1 / 1 ponto

- ☒ Yes
- ☐ No

✓ **Correto**
Correct!

9. **Case 1:** Imagine that you are a vision system. When you are first turned on for the day, an image floods into your camera. You can see lots of things, but not all things. You can't see objects that are occluded, and of course you can't see objects that are behind you. After seeing that first scene, do you have access to the Markov state of the environment?

1 / 1 ponto

Case 2: Imagine that the vision system never worked properly: it always returned the same static image, forever. Would you have access to the Markov state then? (Hint: Reason about $P(S_{t+1} | S_t, \dots, S_0)$, where $S_t = \text{AllWhitePixels}$)

- ☒ You have access to the Markov state in both Case 1 and 2.
- ☐ You have access to the Markov state in Case 1, but you don't have access to the Markov state in Case 2.
- ☐ You don't have access to the Markov state in Case 1, but you do have access to the Markov state in Case 2.
- ☐ You don't have access to the Markov state in both Case 1 and 2.

✓ **Correto**
Correct! Because there is no history before the first image, the first state has the Markov property. The Markov property does not mean that the state representation tells all that would be useful to know, only that it has not forgotten anything that would be useful to know.

The case when the camera is broken is different, but again we have the Markov property. All the possible futures are the same (all white), so nothing needs to be remembered in order to predict them.

10. What is the reward hypothesis?

1 / 1 ponto

- ☐ Ignore rewards and find other signals.
- ☒ Goals and purposes can be thought of as the maximization of the expected value of the cumulative sum of rewards received.
- ☐ Always take the action that gives you the best reward at that point.
- ☐ Goals and purposes can be thought of as the minimization of the expected value of the cumulative sum of rewards received.

✓ **Correto**
Correct!

11. Imagine, an agent is in a maze-like gridworld. You would like the agent to find the goal, as quickly as possible. You give the agent a reward of +1 when it reaches the goal and the discount rate is 1.0, because this is an episodic task. When you run the agent it finds the goal, but does not seem to care how long it takes to complete each episode. How could you fix this? (Select all that apply)

1 / 1 ponto

- ☒ Set a discount rate less than 1 and greater than 0, like 0.9.

✓ **Correto**
Correct! From a given state, the sooner you get the +1 reward, the larger the return. The agent is incentivized to reach the goal faster to maximize expected return.

- ☐ Give the agent a reward of 0 at every time step so it wants to leave.

- ☒ Give the agent -1 at each time step.

✓ **Correto**
Correct! Giving the agent a negative reward on each time step, tells the agent to complete each episode as quickly as possible.

- ☐ Give the agent a reward of +1 at every time step.

12. When may you want to formulate a problem as episodic?

1 / 1 ponto

- ☒ When the agent-environment interaction naturally breaks into sequences. Each sequence begins independently of how the episode ended.
- ☐ When the agent-environment interaction does not naturally break into sequences. Each new episode begins independently of how the previous episode ended.

✓ **Correto**
Correct!